

# Databases and SQL for Data Science with Python

04/08/2024

# CASE Statements

A **CASE** statement in SQL is used to create **conditional logic** within queries. It allows you to perform different actions based on different conditions. Here is the general syntax for a CASE statement in SQL:

```
CASE
  WHEN [first conditional statement]
    THEN [value or calculation]
  WHEN [second conditional statement]
    THEN [value or calculation]
  ELSE [value or calculation]
END
```

This statement indicates that you want a column to contain different values under different conditions.

# CASE Statements

the **WHEN**s are evaluated in order, from **top to bottom**, and the first time a condition evaluates to TRUE, the corresponding THEN part of the statement is executed, and no other WHEN conditions are evaluated. To illustrate, consider this nonsense query:

```
SELECT  
  CASE  
    WHEN 1=1 THEN 'Yes'  
    WHEN 2=2 THEN 'No'  
  END
```

This query will always evaluate to “**Yes**,” because  $1=1$  is always **TRUE**, and therefore the  $2=2$  conditional statement is **never** evaluated, even though it is also **true**.

The **ELSE** part of the statement is **optional**, and that value or calculation result is returned if none of the conditional statements above it evaluate to TRUE.

If the ELSE is not included and none of the WHEN conditionals evaluate to TRUE, the resulting value will be **NULL**.

# CASE Statements

Let's say that we want to know which vendors primarily sell fresh produce and which don't.

| vendor_type                   |
|-------------------------------|
| Arts & Jewelry                |
| Eggs & Meats                  |
| Fresh Focused                 |
| Fresh Variety: Veggies & More |
| Prepared Foods                |

The vendors we want to label as "Fresh Produce" have the word "Fresh" in the `vendor_type` column.

We can use a CASE statement and the LIKE operator to create a new column:

```
SELECT
    vendor_id,
    vendor_name,
    vendor_type,
    CASE
        WHEN LOWER(vendor_type) LIKE '%fresh%'
        THEN 'Fresh Produce'
        ELSE 'Other'
    END AS vendor_type_condensed
FROM farmers_market.vendor
```

# CASE Statements

This gives us:

| vendor_id | vendor_name                      | vendor_type                   | vendor_type_condensed |
|-----------|----------------------------------|-------------------------------|-----------------------|
| 1         | Chris's Sustainable Eggs & Meats | Eggs & Meats                  | Other                 |
| 3         | Hernández Salsa & Veggies        | Fresh Variety: Veggies & More | Fresh Produce         |
| 4         | Mountain View Vegetables         | Fresh Variety: Veggies & More | Fresh Produce         |
| 6         | Seashell Clay Shop               | Arts & Jewelry                | Other                 |
| 7         | Mother's Garlic & Greens         | Fresh Variety: Veggies & More | Fresh Produce         |
| 8         | Marco's Peppers                  | Fresh Focused                 | Fresh Produce         |
| 9         | Annie's Pies                     | Prepared Foods                | Other                 |
| 10        | Mediterranean Bakery             | Prepared Foods                | Other                 |
| 11        | Fields of Corn                   | Fresh Focused                 | Fresh Produce         |

We're using the `LOWER()` function to lowercase the vendor type string, because we don't want the comparison to fail because of capitalization.

# CASE Statements: Creating Binary Flags

A CASE statement can create a binary flag field, commonly found in machine learning datasets.

This field contains only **1s or 0s**, indicating values like "Yes/No" or "Exists/Doesn't exist."

For example, in our database, Farmer's Markets occur on **Wednesday** evenings or **Saturday** mornings. Machine learning algorithms may not understand the "Wednesday" and "Saturday" in the **market\_day** column.

```
SELECT
    market_date,
    market_day
FROM farmers_market.market_date_info
LIMIT 5
```

| market_date | market_day |
|-------------|------------|
| 2019-03-02  | Saturday   |
| 2019-03-09  | Saturday   |
| 2019-03-13  | Wednesday  |
| 2019-03-16  | Saturday   |
| 2019-03-20  | Wednesday  |



# CASE Statements: Creating Binary Flags

We can do this with a CASE statement, making a new column that contains a 1 if the market occurs on a Saturday or Sunday, and a 0 if it doesn't:

```
SELECT
    market_date,
    CASE
        WHEN market_day = 'Saturday' OR market_day = 'Sunday'
        THEN 1
        ELSE 0
    END AS weekend_flag
FROM farmers_market.market_date_info
LIMIT 5
```

| market_date | weekend_flag |
|-------------|--------------|
| 2019-03-02  | 1            |
| 2019-03-09  | 1            |
| 2019-03-13  | 0            |
| 2019-03-16  | 1            |
| 2019-03-20  | 0            |

# CASE Statements: Binning Continuous Values

We initially filtered customer purchases with items costing over \$50 using a conditional statement in the WHERE clause.

Instead, if we want to return all rows and indicate whether the cost was over \$50, we can write the query as follows:

```
SELECT
    market_date,
    customer_id,
    vendor_id,
    ROUND(quantity * cost_to_customer_per_qty, 2) AS price,
    CASE
        WHEN quantity * cost_to_customer_per_qty > 50
        THEN 1
        ELSE 0
    END AS price_over_50
FROM farmers_market.customer_purchases
LIMIT 10
```



# CASE Statements: Binning Continuous Values

The final column now contains a flag indicating which item purchases were over \$50.

| market_date | customer_id | vendor_id | price | price_over_50 |
|-------------|-------------|-----------|-------|---------------|
| 2019-03-02  | 4           | 8         | 8.00  | 0             |
| 2019-03-02  | 10          | 8         | 4.00  | 0             |
| 2019-03-09  | 12          | 8         | 4.00  | 0             |
| 2019-03-09  | 5           | 9         | 16.00 | 0             |
| 2019-03-09  | 1           | 9         | 18.00 | 0             |
| 2019-03-09  | 12          | 9         | 54.00 | 1             |
| 2019-03-02  | 2           | 4         | 9.20  | 0             |
| 2019-03-02  | 3           | 4         | 16.80 | 0             |
| 2019-03-02  | 4           | 4         | 2.80  | 0             |
| 2019-03-09  | 4           | 4         | 19.80 | 0             |

# CASE Statements: Binning Continuous Values

CASE statements can "bin" a continuous variable, like price. For example, we can categorize line-item customer purchases into bins: under \$5.00, \$5.00–\$9.99, \$10.00–\$19.99, and \$20.00 and over.

```
SELECT
    market_date,
    customer_id,
    vendor_id,
    ROUND(quantity * cost_to_customer_per_qty, 2) AS price,
    CASE
        WHEN quantity * cost_to_customer_per_qty < 5.00
            THEN 'Under $5'
        WHEN quantity * cost_to_customer_per_qty < 10.00
            THEN '$5-$9.99'
        WHEN quantity * cost_to_customer_per_qty < 20.00
            THEN '$10-$19.99'
        WHEN quantity * cost_to_customer_per_qty >= 20.00
            THEN '$20 and Up'
        END AS price_bin
    FROM farmers_market.customer_purchases
    LIMIT 10
```

| market_date | customer_id | vendor_id | price | price_bin    |
|-------------|-------------|-----------|-------|--------------|
| 2019-03-02  | 4           | 8         | 8.00  | \$5-\$9.99   |
| 2019-03-02  | 10          | 8         | 4.00  | Under \$5    |
| 2019-03-09  | 12          | 8         | 4.00  | Under \$5    |
| 2019-03-09  | 5           | 9         | 16.00 | \$10-\$19.99 |
| 2019-03-09  | 1           | 9         | 18.00 | \$10-\$19.99 |
| 2019-03-09  | 12          | 9         | 54.00 | \$20 and Up  |
| 2019-03-02  | 2           | 4         | 9.20  | \$5-\$9.99   |
| 2019-03-02  | 3           | 4         | 16.80 | \$10-\$19.99 |
| 2019-03-02  | 4           | 4         | 2.80  | Under \$5    |
| 2019-03-09  | 4           | 4         | 19.80 | \$10-\$19.99 |

# CASE Statements: Binning Continuous Values

Or, if the result needs to be numeric, a different approach is to output the bottom end of the numeric range:

```
SELECT
    market_date,
    customer_id,
    vendor_id,
    ROUND(quantity * cost_to_customer_per_qty, 2) AS price,
    CASE
        WHEN quantity * cost_to_customer_per_qty < 5.00
            THEN 0
        WHEN quantity * cost_to_customer_per_qty < 10.00
            THEN 5
        WHEN quantity * cost_to_customer_per_qty < 20.00
            THEN 10
        WHEN quantity * cost_to_customer_per_qty >= 20.00
            THEN 20
    END AS price_bin_lower_end
FROM farmers_market.customer_purchases
LIMIT 10
```

| market_date | customer_id | vendor_id | price | price_bin_lower_end |
|-------------|-------------|-----------|-------|---------------------|
| 2019-03-02  | 4           | 8         | 8.00  | 5                   |
| 2019-03-02  | 10          | 8         | 4.00  | 0                   |
| 2019-03-09  | 12          | 8         | 4.00  | 0                   |
| 2019-03-09  | 5           | 9         | 16.00 | 10                  |
| 2019-03-09  | 1           | 9         | 18.00 | 10                  |
| 2019-03-09  | 12          | 9         | 54.00 | 20                  |
| 2019-03-02  | 2           | 4         | 9.20  | 5                   |
| 2019-03-02  | 3           | 4         | 16.80 | 10                  |
| 2019-03-02  | 4           | 4         | 2.80  | 0                   |
| 2019-03-09  | 4           | 4         | 19.80 | 10                  |

# CASE Statements: Categorical Encoding

When creating datasets for machine learning, you often need to **encode categorical strings** as **numeric values** for algorithms to use them.

If the categories represent something that can be sorted in a rank order, it might make sense to convert the string variables into numeric values that represent that rank order.

```
SELECT
  booth_number,
  booth_price_level,
  CASE
    WHEN booth_price_level = 'A' THEN 1
    WHEN booth_price_level = 'B' THEN 2
    WHEN booth_price_level = 'C' THEN 3
  END AS booth_price_level_numeric
FROM farmers_market.booth
LIMIT 5
```

| booth_number | booth_price_level | booth_price_level_numeric |
|--------------|-------------------|---------------------------|
| 1            | A                 | 1                         |
| 2            | A                 | 1                         |
| 3            | B                 | 2                         |
| 4            | C                 | 3                         |
| 5            | C                 | 3                         |

# CASE Statements: Categorical Encoding

If the categories aren't necessarily in any kind of rank order, like our vendor type categories, we might use a method called “one-hot encoding.”

This helps us avoid inadvertently indicating a sort order when none exists.

```
SELECT
  vendor_id,
  vendor_name,
  vendor_type,
  CASE WHEN vendor_type = 'Arts & Jewelry'
    THEN 1
    ELSE 0
  END AS vendor_type_arts_jewelry,
  CASE WHEN vendor_type = 'Eggs & Meats'
    THEN 1
    ELSE 0
  END AS vendor_type_eggs_meats,
  CASE WHEN vendor_type = 'Fresh Focused'
    THEN 1
    ELSE 0
  END AS vendor_type_fresh_focused,
  CASE WHEN vendor_type = 'Fresh Variety: Veggies & More'
    THEN 1
    ELSE 0
  END AS vendor_type_fresh_variety,
  CASE WHEN vendor_type = 'Prepared Foods'
    THEN 1
    ELSE 0
  END AS vendor_type_prepared
FROM farmers_market.vendor
```

| vendor_id | vendor_name                      | vendor_type                   | vendor_type_arts_jewelry | vendor_type_eggs_meats | vendor_type_fresh_focused | vendor_type_fresh_variety | vendor_type_prepared |
|-----------|----------------------------------|-------------------------------|--------------------------|------------------------|---------------------------|---------------------------|----------------------|
| 1         | Chris's Sustainable Eggs & Meats | Eggs & Meats                  | 0                        | 1                      | 0                         | 0                         | 0                    |
| 3         | Hernández Salsa & Veggies        | Fresh Variety: Veggies & More | 0                        | 0                      | 0                         | 1                         | 0                    |
| 4         | Mountain View Vegetables         | Fresh Variety: Veggies & More | 0                        | 0                      | 0                         | 1                         | 0                    |
| 6         | Seashell Clay Shop               | Arts & Jewelry                | 1                        | 0                      | 0                         | 0                         | 0                    |
| 7         | Mother's Garlic & Greens         | Fresh Variety: Veggies & More | 0                        | 0                      | 0                         | 1                         | 0                    |
| 8         | Marc's Peppers                   | Fresh Focused                 | 0                        | 0                      | 1                         | 0                         | 0                    |
| 9         | Annie's Pies                     | Prepared Foods                | 0                        | 0                      | 0                         | 0                         | 1                    |
| 10        | Mediterranean Bakery             | Prepared Foods                | 0                        | 0                      | 0                         | 0                         | 1                    |
| 11        | Fields of Corn                   | Fresh Focused                 | 0                        | 0                      | 1                         | 0                         | 0                    |



# CASE Statements: Categorical Encoding

A situation to be aware of when manually encoding one-hot categorical variables this way is that if a new category is added (a new type of vendor in this case), there will be no column in your dataset for the new vendor type until you add another CASE statement.



# Exercise

For the `product` table:

1. Products can be sold by the individual unit or by bulk measures like lbs. or oz. Write a query that outputs the `product_id` and `product_name` columns from the product table, and add a column called `prod_qty_type_condensed` that displays the word "unit" if the `product_qty_type` is "unit," and otherwise displays the word "bulk."
2. We want to flag all of the different types of pepper products that are sold at the market. Add a column to the previous query called `pepper_flag` that outputs a 1 if the `product_name` contains the word "pepper" (regardless of capitalization), and otherwise outputs 0.


# SQL – JOINS

# SQL JOINS

The type of relationship between database tables, and the key fields that connect them, give us information we need to combine them using a **JOIN** statement in SQL.

For example, to list each product with its category name, we need to join the **product** table (which has only category IDs) with the **product\_category** table (which contains category names).

Figure illustrates a **one-to-many** relationship: each product belongs to one category, but each category can include many products. The **primary key** in the **product\_category** table is `product_category_id`, which is also used as a **foreign key** in the **product** table to link products to their categories.



# SQL JOINS

In order to **combine** these tables, we need to figure out which type of **JOIN** to use. To illustrate the different types of SQL JOINS, we'll use the two tables from the Farmer's Market database:



| product_category.<br>product_category_id* | product_category.<br>product_category_name |
|-------------------------------------------|--------------------------------------------|
| 1                                         | Fresh Fruits & Vegetables                  |
| 3                                         | Packaged Prepared Food                     |
| 5                                         | Plants & Flowers                           |
| 6                                         | Eggs & Meat                                |

| product.<br>product_id | product. product_name      | product.<br>product_category_id** |
|------------------------|----------------------------|-----------------------------------|
| 2                      | Jalapeno Peppers - Organic | 1                                 |
| 4                      | Banana Peppers - Jar       | 3                                 |
| 5                      | Whole Wheat Bread          | 3                                 |
| 6                      | Cut Zinnias Bouquet        | 5                                 |
| 7                      | Apple Pie                  | 3                                 |
| 13                     | Baby Salad Lettuce Mix     | 1                                 |
| 99                     | Handmade Candle            | NULL                              |

# SQL JOINS: **LEFT JOIN**

The first type of JOIN we'll cover is the one I've seen used most frequently when building analytical datasets: the **LEFT JOIN**.

This tells the query to pull all records from the table on the "**left side**" of the JOIN, and only the matching records (based on the criteria specified in the JOIN clause) from the table on the "**right side**" of the JOIN.

## Left Join

All rows from the "left table", and only rows from the "right table" with matching values in the specified fields

|  |  | --- |  |
|--|--|-----|--|
|  |  |     |  |
|  |  | --- |  |
|  |  |     |  |
|  |  | --- |  |
|  |  |     |  |
|  |  | --- |  |
|  |  |     |  |
|  |  |     |  |
|  |  |     |  |

# SQL JOINS: LEFT JOIN

In Figure, the row with `product\_id` 99 is included, but the product\_category columns are NULL because there was no matching `product\_category\_id`. The row with `product\_category\_id` 6 is not included because it has no matching record in the left table, as this is a LEFT JOIN.

| product.<br>product_id | product.<br>product_name   | product.<br>product_category_id | product_category.<br>product_category_id | product_category.<br>product_category_name |
|------------------------|----------------------------|---------------------------------|------------------------------------------|--------------------------------------------|
| 2                      | Jalapeno Peppers - Organic | 1                               | 1                                        | Fresh Fruits & Vegetables                  |
| 4                      | Banana Peppers - Jar       | 3                               | 3                                        | Packaged Prepared Food                     |
| 5                      | Whole Wheat Bread          | 3                               | 3                                        | Packaged Prepared Food                     |
| 6                      | Cut Zinnias Bouquet        | 5                               | 5                                        | Plants & Flowers                           |
| 7                      | Apple Pie                  | 3                               | 3                                        | Packaged Prepared Food                     |
| 13                     | Baby Salad Lettuce Mix     | 1                               | 1                                        | Fresh Fruits & Vegetables                  |
| 99                     | Handmade Candle            | NULL                            | NULL                                     | NULL                                       |



# SQL JOINS: **LEFT JOIN**

The syntax for creating this output is:

```
SELECT [columns to return]
FROM [left table]
[JOIN TYPE] [right table]
ON [left table].[field in left table to match] = [right table].[field in right table
to match]
```

# SQL JOINS: **LEFT JOIN**

This is how a query pulling all fields from both tables looks:

```
SELECT * FROM product
      LEFT JOIN product_category
      ON product.product_category_id = product_category.
product_category_id
```

This query can be read as "Select all columns and rows from the product table, and all columns from the product\_category table for rows where the **product\_category's product\_category\_id** matches a product's **product\_category\_id**."

| product_id | product_name               | product_size | product_category_id | product_qty_type | product_category_id | product_category_name         |
|------------|----------------------------|--------------|---------------------|------------------|---------------------|-------------------------------|
| 1          | Habanero Peppers - Organic | medium       | 1                   | lbs              | 1                   | Fresh Fruits & Vegetables     |
| 2          | Jalapeno Peppers - Organic | small        | 1                   | lbs              | 1                   | Fresh Fruits & Vegetables     |
| 3          | Poblano Peppers - Organic  | large        | 1                   | unit             | 1                   | Fresh Fruits & Vegetables     |
| 4          | Banana Peppers - Jar       | 8 oz         | 3                   | unit             | 3                   | Packaged Prepared Food        |
| 5          | Whole Wheat Bread          | 1.5 lbs      | 3                   | unit             | 3                   | Packaged Prepared Food        |
| 6          | Cut Zinnias Bouquet        | medium       | 5                   | unit             | 5                   | Plants & Flowers              |
| 7          | Apple Pie                  | 10"          | 3                   | unit             | 3                   | Packaged Prepared Food        |
| 8          | Cherry Pie                 | 10"          | 3                   | unit             | 3                   | Packaged Prepared Food        |
| 9          | Sweet Potatoes             | medium       | 1                   | lbs              | 1                   | Fresh Fruits & Vegetables     |
| 10         | Eggs                       | 1 dozen      | 6                   | unit             | 6                   | Eggs & Meat (Fresh or Frozen) |

# SQL JOINS: LEFT JOIN

Now, if we want to retrieve specific columns from the merged dataset, we have to specify **which table each column is from**, since it's possible to have identically named columns in different tables. And we can alias identically named columns to differentiate them.

| product_id | product_name                 | product_prod_cat_id | category_prod_cat_id | product_category_name         |
|------------|------------------------------|---------------------|----------------------|-------------------------------|
| 1          | Habanero Peppers - Organic   | 1                   | 1                    | Fresh Fruits & Vegetables     |
| 2          | Jalapeno Peppers - Organic   | 1                   | 1                    | Fresh Fruits & Vegetables     |
| 3          | Poblano Peppers - Organic    | 1                   | 1                    | Fresh Fruits & Vegetables     |
| 4          | Banana Peppers - Jar         | 3                   | 3                    | Packaged Prepared Food        |
| 5          | Whole Wheat Bread            | 3                   | 3                    | Packaged Prepared Food        |
| 6          | Cut Zinnias Bouquet          | 5                   | 5                    | Plants & Flowers              |
| 7          | Apple Pie                    | 3                   | 3                    | Packaged Prepared Food        |
| 8          | Cherry Pie                   | 3                   | 3                    | Packaged Prepared Food        |
| 9          | Sweet Potatoes               | 1                   | 1                    | Fresh Fruits & Vegetables     |
| 10         | Eggs                         | 6                   | 6                    | Eggs & Meat (Fresh or Frozen) |
| 11         | Pork Chops                   | 6                   | 6                    | Eggs & Meat (Fresh or Frozen) |
| 12         | Baby Salad Lettuce Mix - Bag | 1                   | 1                    | Fresh Fruits & Vegetables     |
| 13         | Baby Salad Lettuce Mix       | 1                   | 1                    | Fresh Fruits & Vegetables     |

```
SELECT
    product.product_id,
    product.product_name,
    product.product_category_id AS product_prod_cat_id,
    product_category.product_category_id AS category_prod_cat_id,
    product_category.product_category_name
FROM product
    LEFT JOIN product_category
        ON product.product_category_id = product_category.
product_category_id
```

# SQL JOINS: **LEFT JOIN**

[Table aliasing](#) simplifies SQL queries by allowing you to use short aliases instead of full table names, as shown with aliases "p" and "pc" in the query.

We'll also show only one `product\_category\_id` now that the matching between tables is confirmed.

```
SELECT
    p.product_id,
    p.product_name,
    pc.product_category_id,
    pc.product_category_name
FROM product AS p
    LEFT JOIN product_category AS pc
        ON p.product_category_id = pc.product_category_id
ORDER BY pc.product_category_name, p.product_name
```

| product_id | product_name                 | product_category_id | product_category_name         |
|------------|------------------------------|---------------------|-------------------------------|
| 10         | Eggs                         | 6                   | Eggs & Meat (Fresh or Frozen) |
| 11         | Pork Chops                   | 6                   | Eggs & Meat (Fresh or Frozen) |
| 13         | Baby Salad Lettuce Mix       | 1                   | Fresh Fruits & Vegetables     |
| 12         | Baby Salad Lettuce Mix - Bag | 1                   | Fresh Fruits & Vegetables     |
| 1          | Habanero Peppers - Organic   | 1                   | Fresh Fruits & Vegetables     |
| 2          | Jalapeno Peppers - Organic   | 1                   | Fresh Fruits & Vegetables     |
| 3          | Poblano Peppers - Organic    | 1                   | Fresh Fruits & Vegetables     |
| 9          | Sweet Potatoes               | 1                   | Fresh Fruits & Vegetables     |
| 7          | Apple Pie                    | 3                   | Packaged Prepared Food        |
| 4          | Banana Peppers - Jar         | 3                   | Packaged Prepared Food        |
| 8          | Cherry Pie                   | 3                   | Packaged Prepared Food        |
| 5          | Whole Wheat Bread            | 3                   | Packaged Prepared Food        |
| 6          | Cut Zinnias Bouquet          | 5                   | Plants & Flowers              |

# SQL JOINS: **RIGHT JOIN**

In a **RIGHT JOIN**, all of the rows from the “right table” are returned, along with only the matching rows from the “left table,” using the fields specified in the **ON** part of the query.

## Right Join

All rows from the “right table”, and only rows from the “left table” with matching values in the specified fields

|  |  | --- |  |
|--|--|-----|--|
|  |  | --- |  |
|  |  | --- |  |
|  |  | --- |  |
|  |  | --- |  |
|  |  | --- |  |
|  |  | --- |  |
|  |  | --- |  |



# SQL JOINS: **RIGHT JOIN**

You would use a **RIGHT JOIN** if you wanted to list all product categories and the products in each. (And you didn't care about products that were not put into a category, but you did care about categories that didn't contain any products.)

| product.<br>product_id | product.<br>product_name   | product.<br>product_category_id | product_category.<br>product_category_id | product_category.<br>product_category_name |
|------------------------|----------------------------|---------------------------------|------------------------------------------|--------------------------------------------|
| 2                      | Jalapeno Peppers - Organic | 1                               | 1                                        | Fresh Fruits & Vegetables                  |
| 4                      | Banana Peppers - Jar       | 3                               | 3                                        | Packaged Prepared Food                     |
| 5                      | Whole Wheat Bread          | 3                               | 3                                        | Packaged Prepared Food                     |
| 6                      | Cut Zinnias Bouquet        | 5                               | 5                                        | Plants & Flowers                           |
| 7                      | Apple Pie                  | 3                               | 3                                        | Packaged Prepared Food                     |
| 13                     | Baby Salad Lettuce Mix     | 1                               | 1                                        | Fresh Fruits & Vegetables                  |
| NULL                   | NULL                       | NULL                            | 6                                        | Eggs & Meat                                |



# SQL JOINS: **INNER JOIN**

An **INNER JOIN** only returns records that have matches in both tables.

## Inner Join

Only rows from the “right table” and “left table” where values in the specified fields have matches in both tables

|  |  | --- |  |
|--|--|-----|--|
|  |  |     |  |
|  |  | --- |  |
|  |  |     |  |
|  |  | --- |  |
|  |  | --- |  |
|  |  |     |  |
|  |  |     |  |

# SQL JOINS: **INNER JOIN**

In the **INNER JOIN**ed output, there is a row for every matched pair of **product\_category\_ids**, and no rows without a matching **product\_category\_id** on the other side.

| product.<br>product_id | product.<br>product_name   | product.<br>product_category_id |  | product_category.<br>product_category_id | product_category.<br>product_category_name |
|------------------------|----------------------------|---------------------------------|--|------------------------------------------|--------------------------------------------|
| 2                      | Jalapeno Peppers - Organic | 1                               |  | 1                                        | Fresh Fruits & Vegetables                  |
| 4                      | Banana Peppers - Jar       | 3                               |  | 3                                        | Packaged Prepared Food                     |
| 5                      | Whole Wheat Bread          | 3                               |  | 3                                        | Packaged Prepared Food                     |
| 6                      | Cut Zinnias Bouquet        | 5                               |  | 5                                        | Plants & Flowers                           |
| 7                      | Apple Pie                  | 3                               |  | 3                                        | Packaged Prepared Food                     |
| 13                     | Baby Salad Lettuce Mix     | 1                               |  | 1                                        | Fresh Fruits & Vegetables                  |

# SQL JOINS

To practice all of these types of JOINS, let's now look at the `customer` and `customer_purchase` tables in the Farmer's Market database. Again, this is a **one-to-many** type relationship.

We do a **LEFT JOIN**, using the following query:

```
SELECT *
FROM customer AS c
LEFT JOIN customer_purchases AS cp
ON c.customer_id = cp.customer_id
```

| customer_id | customer_first_name | customer_last_name | customer_zip | product_id | vendor_id | market_date | customer_id | quantity | cost_to_cust | transacti |
|-------------|---------------------|--------------------|--------------|------------|-----------|-------------|-------------|----------|--------------|-----------|
| 5           | Abigail             | Harris             | 22801        | 7          | 9         | 2019-03-09  | 5           | 1.00     | 16.00        | 10:41:00  |
| 6           | Betty               | Bullard            | 22801        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |
| 7           | Jessica             | Armenta            | 22803        | 12         | 7         | 2019-03-09  | 7           | 2.00     | 3.00         | 11:40:00  |
| 8           | Norma               | Valenzuela         | 22803        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |
| 9           | Janet               | Forbes             | 22801        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |
| 10          | Russell             | Edwards            | 22801        | 4          | 8         | 2019-03-02  | 10          | 1.00     | 4.00         | 09:12:00  |
| 11          | Richard             | Paulson            | 22801        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |
| 12          | Jack                | Wise               | 22803        | 4          | 8         | 2019-03-09  | 12          | 1.00     | 4.00         | 13:03:00  |
| 12          | Jack                | Wise               | 22803        | 8          | 9         | 2019-03-09  | 12          | 3.00     | 18.00        | 13:18:00  |
| 12          | Jack                | Wise               | 22803        | 11         | 1         | 2019-03-09  | 12          | 0.90     | 12.00        | 13:10:00  |
| 12          | Jack                | Wise               | 22803        | 12         | 7         | 2019-03-09  | 12          | 2.00     | 3.00         | 13:00:00  |
| 13          | Jeremy              | Gruber             | 22803        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |
| 14          | William             | Lones              | 22801        | NULL       | NULL      | NULL        | NULL        | NULL     | NULL         | NULL      |

# SQL JOINS

Unlike the `product-product_category` relationship, some customers may not have made any purchases. These customers, who signed up for the loyalty card but haven't bought anything yet, will appear in the output due to the LEFT JOIN.

The list includes all customers with their purchases if available. Customers with multiple purchases will appear multiple times, while those with no purchases will show NULLs for fields from the `customer_purchases` table.

We can use the WHERE clause to filter the list to only customers with no purchases, if we'd like:

```
SELECT c.*  
FROM customer AS c  
LEFT JOIN customer_purchases AS cp  
    ON c.customer_id = cp.customer_id  
WHERE cp.customer_id IS NULL
```

| customer_id | customer_first_name | customer_last_name | customer_zip |
|-------------|---------------------|--------------------|--------------|
| 6           | Betty               | Bullard            | 22801        |
| 8           | Norma               | Valenzuela         | 22803        |
| 9           | Janet               | Forbes             | 22801        |
| 11          | Richard             | Paulson            | 22801        |
| 13          | Jeremy              | Gruber             | 22803        |
| 14          | William             | Lopes              | 22801        |
| 15          | Darrell             | Messina            | 22801        |

# SQL JOINS

To list all purchases and their associated customers, we would use a RIGHT JOIN. This will pull all records from the `customer_purchases` table and only include customers from the customer table who have made a purchase.

```
SELECT *
FROM customer AS c
RIGHT JOIN customer_purchases AS cp
ON c.customer_id = cp.customer_id
```

there are no rows returned with `NULL` values in the customer table columns, because there is no such thing as a purchase without a `customer_id`.

| customer_id | customer_first_name | customer_last_name | customer_zip | product_id | vendor_id | market_date | customer_id | quantity | cost_to_cust | transacti |
|-------------|---------------------|--------------------|--------------|------------|-----------|-------------|-------------|----------|--------------|-----------|
| 4           | Deanna              | Washington         | 22801        | 4          | 8         | 2019-03-02  | 4           | 2.00     | 4.00         | 10:22:00  |
| 10          | Russell             | Edwards            | 22801        | 4          | 8         | 2019-03-02  | 10          | 1.00     | 4.00         | 09:12:00  |
| 12          | Jack                | Wise               | 22803        | 4          | 8         | 2019-03-09  | 12          | 1.00     | 4.00         | 13:03:00  |
| 5           | Abigail             | Harris             | 22801        | 7          | 9         | 2019-03-09  | 5           | 1.00     | 16.00        | 10:41:00  |
| 1           | Jane                | Connor             | 22801        | 8          | 9         | 2019-03-09  | 1           | 1.00     | 18.00        | 08:25:00  |
| 12          | Jack                | Wise               | 22803        | 8          | 9         | 2019-03-09  | 12          | 3.00     | 18.00        | 13:18:00  |
| 2           | Manuel              | Diaz               | 22803        | 9          | 4         | 2019-03-02  | 2           | 4.60     | 2.00         | 10:53:00  |
| 3           | Bob                 | Wilson             | 22803        | 9          | 4         | 2019-03-02  | 3           | 8.40     | 2.00         | 11:39:00  |
| 4           | Deanna              | Washington         | 22801        | 9          | 4         | 2019-03-02  | 4           | 1.40     | 2.00         | 10:31:00  |
| 4           | Deanna              | Washington         | 22801        | 9          | 4         | 2019-03-09  | 4           | 9.90     | 2.00         | 13:02:00  |
| 1           | Jane                | Connor             | 22801        | 10         | 1         | 2019-03-02  | 1           | 1.00     | 5.50         | 08:59:00  |
| 1           | Jane                | Connor             | 22801        | 10         | 1         | 2019-03-02  | 1           | 3.00     | 5.00         | 09:31:00  |
| 1           | Jane                | Connor             | 22801        | 10         | 1         | 2019-03-09  | 1           | 2.00     | 5.50         | 08:30:00  |

# SQL JOINS

If you only want records from each table that have matches in both tables, use an **INNER JOIN**.

Using these customer and **customer\_purchases** tables, an INNER JOIN happens to return the same results as the **RIGHT JOIN**, because there aren't any records on the "**right side**" of the join without matches on the "**left side**"—every purchase is associated with a customer.



# SQL JOINS: Pitfall when Filtering Joined Data

How do you think the output of the following query will differ from the original **LEFT JOIN** query without the added **WHERE** clause?

```
SELECT *  
FROM customer AS c  
LEFT JOIN customer_purchases AS cp  
      ON c.customer_id = cp.customer_id  
WHERE cp.customer_id > 0
```

Although `customer_id` values are all integers above 0, adding this **WHERE** clause filters based on the `customer_id` in the `customer_purchases` table (alias `cp`). This removes customers without purchases, making the query behave like an **INNER JOIN** by excluding records with **NULL** values in the **right-side** table columns.

When using a **LEFT JOIN** to return all rows from the **left** table, avoid filtering on fields from the right table unless you allow **NULLs**. Otherwise, you might filter out rows you intended to keep.

# SQL JOINS: Pitfall when Filtering Joined Data

Let's say we want to write a query that returns a list of all customers who did not make a purchase at the March 2, 2019, farmer's market.

We will use a **LEFT JOIN**, since we want to include the customers who have never made a purchase at any farmer's market, so wouldn't have any records in the **customer\_purchases** table:

```
SELECT c.*, cp.market_date
FROM customer AS c
LEFT JOIN customer_purchases AS cp
      ON c.customer_id = cp.customer_id
WHERE cp.market_date <> '2019-03-02'
```

| customer_id | customer_first_name | customer_last_name | customer_zip | market_date |
|-------------|---------------------|--------------------|--------------|-------------|
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 2           | Manuel              | Diaz               | 22803        | 2019-03-13  |
| 2           | Manuel              | Diaz               | 22803        | 2019-03-13  |
| 3           | Bob                 | Wilson             | 22803        | 2019-03-16  |
| 3           | Bob                 | Wilson             | 22803        | 2019-03-16  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 5           | Abigail             | Harris             | 22801        | 2019-03-09  |
| 7           | Jessica             | Armenta            | 22803        | 2019-03-09  |
| 10          | Russell             | Edwards            | 22801        | 2019-03-16  |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |
| 12          | Jack                | Wise               | 22803        | 2019-03-16  |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |

# SQL JOINS: Pitfall when Filtering Joined Data

There are multiple problems with this output:

1. The problem is that we're missing customers who never made a purchase because we filtered on `market_date` from the `customer_purchases` table (the right side of the JOIN).

Since SQL doesn't compare NULL values to TRUE, this filter excludes customers without purchases, although we need it to exclude those who made a purchase on that day.

**Solution:** To filter results using a field from the right table while still returning records from the left table, adjust the WHERE clause to allow NULL values.

```
SELECT c.*, cp.market_date
FROM customer AS c
LEFT JOIN customer_purchases AS cp
      ON c.customer_id = cp.customer_id
WHERE (cp.market_date <> '2019-03-02' OR cp.market_date IS NULL)
```

# SQL JOINS: Pitfall when Filtering Joined Data

Now we see customers without purchases:

| customer_id | customer_first_name | customer_last_name | customer_zip | market_date |
|-------------|---------------------|--------------------|--------------|-------------|
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 1           | Jane                | Connor             | 22801        | 2019-03-09  |
| 2           | Manuel              | Diaz               | 22803        | 2019-03-13  |
| 2           | Manuel              | Diaz               | 22803        | 2019-03-13  |
| 3           | Bob                 | Wilson             | 22803        | 2019-03-16  |
| 3           | Bob                 | Wilson             | 22803        | 2019-03-16  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 4           | Deanna              | Washington         | 22801        | 2019-03-09  |
| 5           | Abigail             | Harris             | 22801        | 2019-03-09  |
| 6           | Betty               | Bullard            | 22801        | NULL        |
| 7           | Jessica             | Armenta            | 22803        | 2019-03-09  |
| 8           | Norma               | Valenzuela         | 22803        | NULL        |
| 9           | Janet               | Forbes             | 22801        | NULL        |
| 10          | Russell             | Edwards            | 22801        | 2019-03-16  |
| 11          | Richard             | Paulson            | 22801        | NULL        |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |
| 12          | Jack                | Wise               | 22803        | 2019-03-09  |
| 12          | Jack                | Wise               | 22803        | 2019-03-16  |

# SQL JOINS: Pitfall when Filtering Joined Data

There are multiple problems with this output:

2. The output shows one row per customer per item purchased, rather than just a list of customers, because the `customer\_purchases` table includes a record for each purchased item.

**Solution:** Removing the `market\_date` field from the `customer\_purchases` table and using the **DISTINCT** keyword to display only unique results.

```
SELECT DISTINCT c.*  
FROM customer AS c  
LEFT JOIN customer_purchases AS cp  
    ON c.customer_id = cp.customer_id  
WHERE (cp.market_date <> '2019-03-02' OR cp.market_date IS NULL)
```

# SQL JOINS: Pitfall when Filtering Joined Data

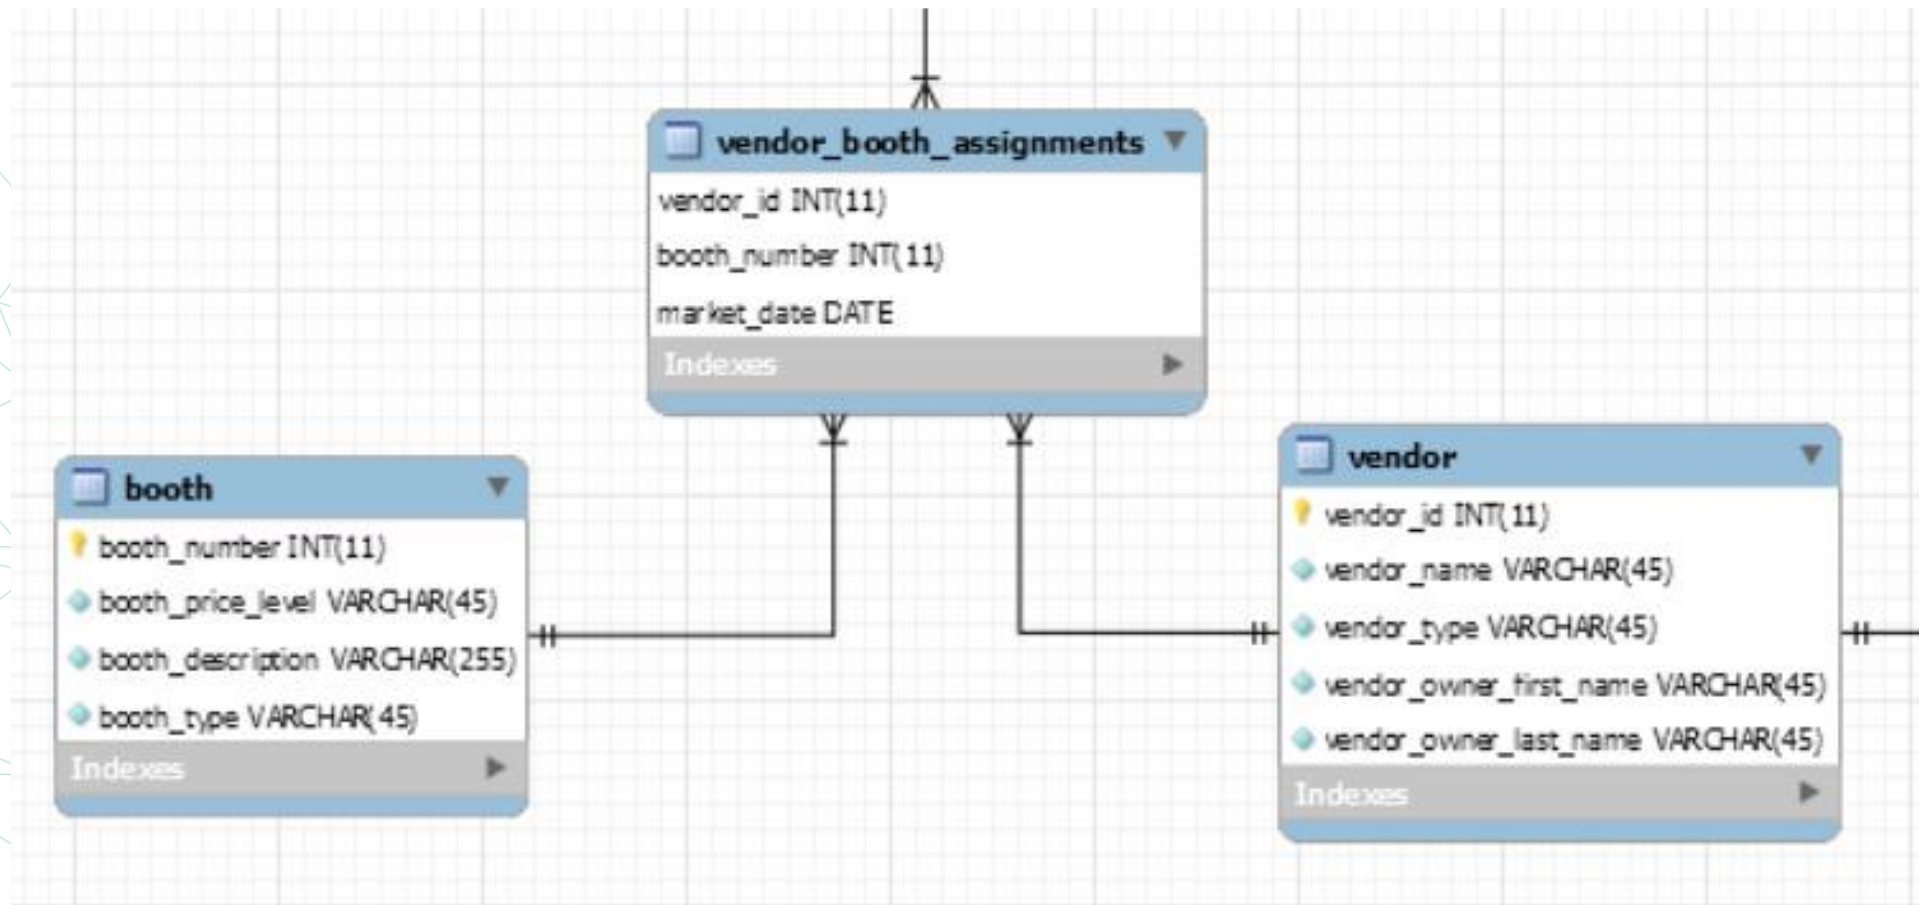
This results:

| customer_id | customer_first_name | customer_last_name | customer_zip |
|-------------|---------------------|--------------------|--------------|
| 1           | Jane                | Connor             | 22801        |
| 2           | Manuel              | Díaz               | 22803        |
| 3           | Bob                 | Wilson             | 22803        |
| 4           | Deanna              | Washington         | 22801        |
| 5           | Abigail             | Harris             | 22801        |
| 6           | Betty               | Bullard            | 22801        |
| 7           | Jessica             | Armenta            | 22803        |
| 8           | Norma               | Valenzuela         | 22803        |
| 9           | Janet               | Forbes             | 22801        |
| 10          | Russell             | Edwards            | 22801        |
| 11          | Richard             | Paulson            | 22801        |
| 12          | Jack                | Wise               | 22803        |
| 13          | Jeremy              | Gruber             | 22803        |
| 14          | William             | Lopes              | 22801        |
| 15          | Darrell             | Messina            | 22801        |



# SQL JOINS: More than Two Tables

To get details about all farmer's market booths and vendor assignments for every market date, we need to join the three tables to create a merged dataset:



# SQL JOINS: More than Two Tables

What JOINS can we use to include all booths, even those not assigned to a vendor, and all vendors assigned to booths?

We can **LEFT JOIN** the **vendor\_booth\_assignments** to **booth**, therefore including all of the booths, and **LEFT JOIN** **vendor** to **vendor\_booth\_assignments** in the results.

```
SELECT
    b.booth_number,
    b.booth_type,
    vba.market_date,
    v.vendor_id,
    v.vendor_name,
    v.vendor_type
FROM booth AS b
    LEFT JOIN vendor_booth_assignments AS vba ON b.booth_number = vba.
booth_number
    LEFT JOIN vendor AS v ON v.vendor_id = vba.vendor_id
ORDER BY b.booth_number, vba.market_date
```

# SQL JOINS: More than Two Tables

This results:

| booth_number | booth_type | market_date | vendor_id | vendor_name                      | vendor_type                   |
|--------------|------------|-------------|-----------|----------------------------------|-------------------------------|
| 1            | Standard   | 2019-03-02  | 3         | Hernández Salsa & Veggies        | Fresh Variety: Veggies & More |
| 1            | Standard   | 2019-03-09  | 3         | Hernández Salsa & Veggies        | Fresh Variety: Veggies & More |
| 1            | Standard   | 2019-03-13  | 3         | Hernández Salsa & Veggies        | Fresh Variety: Veggies & More |
| 2            | Standard   | 2019-03-02  | 1         | Chris's Sustainable Eggs & Meats | Eggs & Meats                  |
| 2            | Standard   | 2019-03-09  | 1         | Chris's Sustainable Eggs & Meats | Eggs & Meats                  |
| 2            | Standard   | 2019-03-13  | 4         | Mountain View Vegetables         | Fresh Variety: Veggies & More |
| 3            | Small      | NULL        | NULL      | NULL                             | NULL                          |
| 4            | Small      | NULL        | NULL      | NULL                             | NULL                          |
| 5            | Small      | NULL        | NULL      | NULL                             | NULL                          |
| 6            | Small      | 2019-03-02  | 8         | Marco's Peppers                  | Fresh Focused                 |
| 6            | Small      | 2019-03-09  | 8         | Marco's Peppers                  | Fresh Focused                 |
| 6            | Small      | 2019-03-13  | 8         | Marco's Peppers                  | Fresh Focused                 |
| 7            | Standard   | 2019-03-02  | 4         | Mountain View Vegetables         | Fresh Variety: Veggies & More |
| 7            | Standard   | 2019-03-09  | 4         | Mountain View Vegetables         | Fresh Variety: Veggies & More |
| 8            | Small      | 2019-03-02  | 9         | Annie's Pies                     | Prepared Foods                |
| 8            | Small      | 2019-03-09  | 9         | Annie's Pies                     | Prepared Foods                |
| 8            | Small      | 2019-03-13  | 10        | Mediterranean Bakery             | Prepared Foods                |

# SQL JOINS: More than Two Tables

You can think of the second JOIN as being merged into the result of the first JOIN.

## Table LEFT JOINed to a table on the RIGHT side of an existing LEFT JOIN

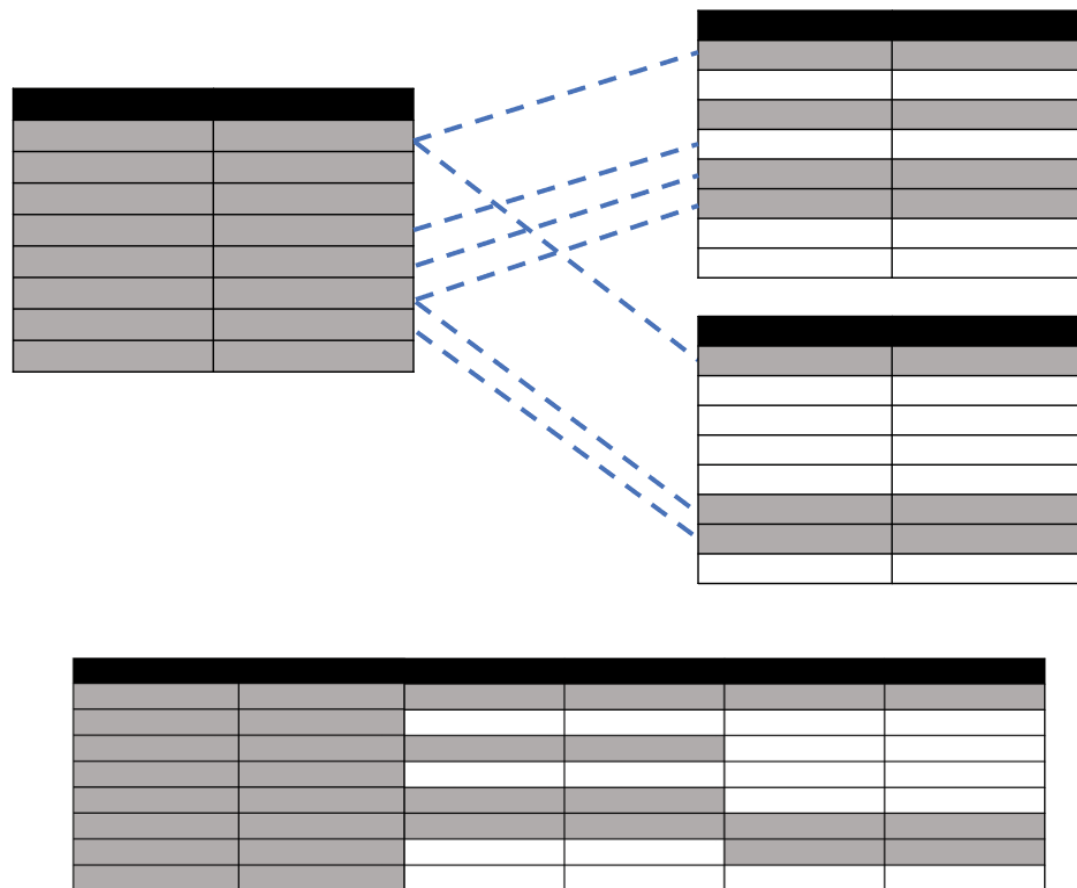
All rows from the "left table", only rows from the "middle table" with matching values in the specified fields of the "left table", and only rows from the "right table" with matching values in the specified fields of the "middle table".

# SQL JOINS: More than Two Tables

## Two Tables LEFT JOINed to a Table

All rows from the "left table", and only rows from each "right table" with matching values in the specified fields of the "left table".

If the third table joined the first table using a common field (though not possible with the Farmer's Market database since no other tables join the booth table), the arrangement would look like this diagram.



# Exercise

1. Write a query that **INNER JOINs** the **vendor** table to the **vendor\_booth\_assignments** table on the **vendor\_id** field they both have in common, and sorts the result by **vendor\_name**, then **market\_date**.
2. Is it possible to write a query that produces an output identical to the output of the following query, but using a **LEFT JOIN** instead of a **RIGHT JOIN**?

```
SELECT *  
FROM customer AS c  
RIGHT JOIN customer_purchases AS cp  
ON c.customer_id = cp.customer_id
```



# SQL – Aggregating Results for Analysis

# Aggregation - **GROUP BY** Syntax

We saw this basic SQL SELECT query syntax, But two sections of this query that we haven't yet covered, which are both related to [aggregation](#), are the [GROUP BY](#) and [HAVING](#) clauses:

```
SELECT [columns to return]
FROM [table]
WHERE [conditional filter statements]
GROUP BY [columns to group on]
HAVING [conditional filter statements that are run after grouping]
ORDER BY [columns to sort on]
```

# Aggregation - **GROUP BY** Syntax

Using what you've learned so far without grouping, you might write a query like the following to get a list of the customer IDs of customers who made purchases on each market date:

```
SELECT
    market_date,
    customer_id
FROM farmers_market.customer_purchases
ORDER BY market_date, customer_id
```

However, this approach would result in one row per item each customer purchased, displaying duplicates in the output.

# Aggregation - **GROUP BY** Syntax

To instead get one row per customer per market date, you can group the results by adding a **GROUP BY** clause that specifies that you want to summarize the results by the `customer_id` and `market_date` fields:

```
SELECT
    market_date,
    customer_id
FROM farmers_market.customer_purchases
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
```

# Aggregation - Displaying Group Summaries

Now that you have grouped the data at the desired level, you can add [aggregate](#) functions like [SUM](#) and [COUNT](#) to return summaries of the [customer\\_purchases](#) data per group.

```
SELECT
    market_date,
    customer_id,
    COUNT(*) AS items_purchased
FROM farmers_market.customer_purchases
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
LIMIT 10
```

| market_date | customer_id | items_purchased |
|-------------|-------------|-----------------|
| 2019-03-02  | 1           | 3               |
| 2019-03-02  | 2           | 1               |
| 2019-03-02  | 3           | 1               |
| 2019-03-02  | 4           | 2               |
| 2019-03-02  | 10          | 1               |
| 2019-03-09  | 1           | 4               |
| 2019-03-09  | 4           | 5               |
| 2019-03-09  | 5           | 1               |
| 2019-03-09  | 7           | 1               |
| 2019-03-09  | 12          | 4               |

# Aggregation - Displaying Group Summaries

Remember that the granularity of the `customer_purchases` table means that if a customer buys three identical items at once, it appears as one row with a quantity of 3. However, if they make separate purchases, it generates multiple rows.

To count all quantities purchased rather than just line items, sum the quantity column with the following query.

```
SELECT
    market_date,
    customer_id,
    SUM(quantity) AS items_purchased
FROM farmers_market.customer_purchases
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
LIMIT 10
```



# Aggregation - Displaying Group Summaries

The `items_purchased` column is no longer an integer, because some of the quantities we're adding up are bulk product weights.

| market_date | customer_id | items_purchased |
|-------------|-------------|-----------------|
| 2019-03-02  | 1           | 5.70            |
| 2019-03-02  | 2           | 4.60            |
| 2019-03-02  | 3           | 8.40            |
| 2019-03-02  | 4           | 3.40            |
| 2019-03-02  | 10          | 1.00            |
| 2019-03-09  | 1           | 5.20            |
| 2019-03-09  | 4           | 13.20           |
| 2019-03-09  | 5           | 1.00            |
| 2019-03-09  | 7           | 2.00            |
| 2019-03-09  | 12          | 6.90            |

# Aggregation - Displaying Group Summaries

After seeing these results and realizing bulk quantities are included, you may decide it makes more sense to know how many different kinds of items each customer purchased.

So, you only want to count "1" if they bought tomatoes, regardless of the quantity or number of purchases, and add to that count only if they bought other items.

What you want now is a **DISTINCT** count of product IDs, shown in the following query:

```
SELECT
    market_date,
    customer_id,
    COUNT(DISTINCT product_id) AS different_products_purchased
FROM farmers_market.customer_purchases c
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
LIMIT 10
```

# Aggregation - Displaying Group Summaries

Now we're identifying how many unique `product_id` values exist across those rows in the group—how many different kinds of products were purchased by each customer on each market date:

| market_date | customer_id | different_products_purchased |
|-------------|-------------|------------------------------|
| 2019-03-02  | 1           | 2                            |
| 2019-03-02  | 2           | 1                            |
| 2019-03-02  | 3           | 1                            |
| 2019-03-02  | 4           | 2                            |
| 2019-03-02  | 10          | 1                            |
| 2019-03-09  | 1           | 3                            |
| 2019-03-09  | 4           | 4                            |
| 2019-03-09  | 5           | 1                            |
| 2019-03-09  | 7           | 1                            |
| 2019-03-09  | 12          | 4                            |

# Aggregation - Displaying Group Summaries

We can also combine these summaries into a single query:

```
SELECT
    market_date,
    customer_id,
    SUM(quantity) AS items_purchased,
    COUNT(DISTINCT product_id) AS different_products_purchased
FROM farmers_market.customer_purchases
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
LIMIT 10
```

| market_date | customer_id | items_purchased | different_products_purchased |
|-------------|-------------|-----------------|------------------------------|
| 2019-03-02  | 1           | 5.70            | 2                            |
| 2019-03-02  | 2           | 4.60            | 1                            |
| 2019-03-02  | 3           | 8.40            | 1                            |
| 2019-03-02  | 4           | 3.40            | 2                            |
| 2019-03-02  | 10          | 1.00            | 1                            |
| 2019-03-09  | 1           | 5.20            | 3                            |
| 2019-03-09  | 4           | 13.20           | 4                            |
| 2019-03-09  | 5           | 1.00            | 1                            |
| 2019-03-09  | 7           | 2.00            | 1                            |
| 2019-03-09  | 12          | 6.90            | 4                            |

# Aggregation - Calculations Inside

You can also include [mathematical](#) operations, which are calculated at the row level prior to summarization, inside the aggregate functions.

You learned to display a list of customer purchases at the farmer's market using a [WHERE](#) clause to filter for a specific customer.

```
SELECT
    market_date,
    customer_id,
    vendor_id,
    quantity * cost_to_customer_per_qty AS price
FROM farmers_market.customer_purchases
WHERE
    customer_id = 3
ORDER BY market_date, vendor_id
```

| market_date | customer_id | vendor_id | price   |
|-------------|-------------|-----------|---------|
| 2019-03-02  | 3           | 4         | 16.8000 |
| 2019-03-16  | 3           | 4         | 11.0000 |
| 2019-03-16  | 3           | 9         | 18.0000 |



# Aggregation - Calculations Inside

Let's say we wanted to know how much money this customer spent total on each market\_date, regardless of item or vendor.

We can `GROUP BY market_date`, and use the `SUM` aggregate function on the price calculation to add up the prices of the items purchased:

```
SELECT
    customer_id,
    market_date,
    SUM(quantity * cost_to_customer_per_qty) AS total_spent
FROM farmers_market.customer_purchases
WHERE
    customer_id = 3
GROUP BY market_date
ORDER BY market_date
```



# Aggregation - Calculations Inside

What if we wanted to find out how much this customer had spent at each vendor, regardless of date? Then we can group by `customer_id` and `vendor_id`:

```
SELECT
    customer_id,
    vendor_id,
    SUM(quantity * cost_to_customer_per_qty) AS total_spent
FROM farmers_market.customer_purchases
WHERE
    customer_id = 3
GROUP BY customer_id, vendor_id
ORDER BY customer_id, vendor_id
```

| customer_id | vendor_id | total_spent |
|-------------|-----------|-------------|
| 3           | 4         | 27.8000     |
| 3           | 9         | 18.0000     |

# Aggregation - Calculations Inside

We can also remove the `customer_id` filter—in this case by removing the entire WHERE clause —and `GROUP BY customer_id` only, to get a list of every customer and how much they have ever spent at the farmer's market.

```
SELECT
    customer_id,
    SUM(quantity * cost_to_customer_per_qty) AS total_spent
FROM farmers_market.customer_purchases
GROUP BY customer_id
ORDER BY customer_id
```

| customer_id | total_spent |
|-------------|-------------|
| 1           | 100.3750    |
| 2           | 25.4000     |
| 3           | 45.8000     |
| 4           | 66.6250     |
| 5           | 16.0000     |
| 7           | 15.0000     |
| 10          | 8.0000      |
| 12          | 95.8000     |

# Aggregation - Calculations Inside

So far, we have been doing all of this aggregation on a single table, but it can be done on [joined tables](#), as well.

It's a good idea to join the tables without the aggregate functions first, to make sure the data is at the level of granularity you expect (and not generating duplicates) before adding the GROUP BY.

Let's say that for the query that was grouped by [customer\\_id](#) and [vendor\\_id](#), we want to bring in some customer details, such as first and last name, and the vendor name.

# Aggregation - Calculations Inside

We can first join the three tables together, select columns from all of the tables, and inspect the output before grouping.

```
SELECT
    c.customer_first_name,
    c.customer_last_name,
    cp.customer_id,
    v.vendor_name,
    cp.vendor_id,
    cp.quantity * cp.cost_to_customer_per_qty AS price
FROM farmers_market.customer c
    LEFT JOIN farmers_market.customer_purchases cp
        ON c.customer_id = cp.customer_id
    LEFT JOIN farmers_market.vendor v
        ON cp.vendor_id = v.vendor_id
WHERE
    cp.customer_id = 3
ORDER BY cp.customer_id, cp.vendor_id
```

| customer_first_name | customer_last_name | customer_id | vendor_name              | vendor_id | price   |
|---------------------|--------------------|-------------|--------------------------|-----------|---------|
| Bob                 | Wilson             | 3           | Mountain View Vegetables | 4         | 16.8000 |
| Bob                 | Wilson             | 3           | Mountain View Vegetables | 4         | 11.0000 |
| Bob                 | Wilson             | 3           | Annie's Pies             | 9         | 18.0000 |

# Aggregation - Calculations Inside

To summarize **with one row per customer per vendor**, group by all fields from both the **customer** and **vendor** tables that aren't aggregate functions.

```
SELECT
    c.customer_first_name,
    c.customer_last_name,
    cp.customer_id,
    v.vendor_name,
    cp.vendor_id,
    ROUND(SUM(quantity * cost_to_customer_per_qty), 2) AS total_spent
FROM farmers_market.customer c
    LEFT JOIN farmers_market.customer_purchases cp
        ON c.customer_id = cp.customer_id
    LEFT JOIN farmers_market.vendor v
        ON cp.vendor_id = v.vendor_id
WHERE
    cp.customer_id = 3
GROUP BY
    c.customer_first_name,
    c.customer_last_name,
    cp.customer_id,
    v.vendor_name,
    cp.vendor_id
ORDER BY cp.customer_id, cp.vendor_id
```

| customer_first_name | customer_last_name | customer_id | vendor_name              | vendor_id | total_spent |
|---------------------|--------------------|-------------|--------------------------|-----------|-------------|
| Bob                 | Wilson             | 3           | Mountain View Vegetables | 4         | 27.80       |
| Bob                 | Wilson             | 3           | Annie's Pies             | 9         | 18.00       |

# Aggregation - Calculations Inside

We can keep the same aggregation and filter to a single vendor instead of a single customer, to get a list of customers per vendor instead of vendors per customer

```
SELECT
    c.customer_first_name,
    c.customer_last_name,
    cp.customer_id,
    v.vendor_name,
    cp.vendor_id,
    ROUND(SUM(quantity * cost_to_customer_per_qty), 2) AS total_spent
FROM farmers_market.customer c
    LEFT JOIN farmers_market.customer_purchases cp
        ON c.customer_id = cp.customer_id
    LEFT JOIN farmers_market.vendor v
        ON cp.vendor_id = v.vendor_id
WHERE
    cp.vendor_id = 9
GROUP BY
    c.customer_first_name,
    c.customer_last_name,
    cp.customer_id,
    v.vendor_name,
    cp.vendor_id
ORDER BY cp.customer_id, cp.vendor_id
```

Or, we could remove the **WHERE** clause altogether and get one row for every **customer-vendor pair** in the database.

| customer_first_name | customer_last_name | customer_id | vendor_name  | vendor_id | total_spent |
|---------------------|--------------------|-------------|--------------|-----------|-------------|
| Jane                | Connor             | 1           | Annie's Pies | 9         | 18.00       |
| Bob                 | Wilson             | 3           | Annie's Pies | 9         | 18.00       |
| Abigail             | Harris             | 5           | Annie's Pies | 9         | 16.00       |
| Jack                | Wise               | 12          | Annie's Pies | 9         | 72.00       |



# Aggregation – MIN and MAX

To get the **most** and **least** expensive items per product category, use the **vendor\_inventory** table, which lists the original prices set by vendors.

This considers that vendors can adjust prices per customer, as reflected in the **cost\_to\_customer\_per\_qty** field in the **customer\_purchases** table.

First, let's look at all of the available fields in the **vendor\_inventory** table:

```
SELECT *  
FROM farmers_market.vendor_inventory  
ORDER BY original_price  
LIMIT 10
```

| market_date | quantity | vendor_id | product_id | original_price |
|-------------|----------|-----------|------------|----------------|
| 2019-03-09  | 10.00    | 9         | 5          | 5.00           |
| 2019-03-30  | 17.00    | 7         | 13         | 6.00           |
| 2019-03-23  | 8.00     | 7         | 13         | 6.00           |
| 2019-03-02  | 13.00    | 1         | 10         | 6.00           |
| 2019-03-09  | 17.00    | 1         | 10         | 6.00           |
| 2019-03-09  | 8.00     | 7         | 13         | 6.00           |
| 2019-03-20  | 13.00    | 7         | 13         | 6.00           |
| 2019-03-02  | 28.00    | 1         | 11         | 12.00          |
| 2019-03-09  | 10.00    | 1         | 11         | 12.00          |
| 2019-03-20  | 15.00    | 1         | 11         | 13.00          |

# Aggregation - MIN and MAX

We can get the least and most expensive item prices in the entire table by using the `MIN()` and `MAX()` functions without grouping in MySQL:

```
SELECT  
    MIN(original_price) AS minimum_price,  
    MAX(original_price) AS maximum_price  
FROM farmers_market.vendor_inventory  
ORDER BY original_price
```

| minimum_price | maximum_price |
|---------------|---------------|
| 2.00          | 18.00         |

# Aggregation - MIN and MAX

But if we want to get the **lowest and highest prices within each product category**, we have to group by the **product\_category\_id** (and **product\_category\_name**, if we want to display it), then the summary values will be calculated per group:

```
SELECT
    pc.product_category_name,
    p.product_category_id,
    MIN(vi.original_price) AS minimum_price,
    MAX(vi.original_price) AS maximum_price
FROM farmers_market.vendor_inventory AS vi
    INNER JOIN farmers_market.product AS p
        ON vi.product_id = p.product_id
    INNER JOIN farmers_market.product_category AS pc
        ON p.product_category_id = pc.product_category_id
GROUP BY pc.product_category_name, p.product_category_id
```

| product_category_name         | product_category_id | minimum_price | maximum_price |
|-------------------------------|---------------------|---------------|---------------|
| Fresh Fruits & Vegetables     | 1                   | 2.00          | 6.00          |
| Packaged Prepared Food        | 3                   | 4.00          | 18.00         |
| Eggs & Meat (Fresh or Frozen) | 6                   | 6.00          | 13.00         |

# Aggregation - COUNT and COUNT DISTINCT

Suppose we wanted to count how many products were for sale on each market date, or how many different products each vendor offered. We can determine these values using **COUNT** and **COUNT DISTINCT**.

**COUNT** will count up the rows within a group when used with **GROUP BY**, and **COUNT DISTINCT** will count up the unique values present in the specified field within the group.

```
SELECT
    market_date,
    COUNT(product_id) AS product_count
FROM farmers_market.vendor_inventory
GROUP BY market_date
ORDER BY market_date
```

| market_date | product_count |
|-------------|---------------|
| 2019-03-02  | 4             |
| 2019-03-09  | 9             |
| 2019-03-13  | 2             |
| 2019-03-16  | 3             |
| 2019-03-20  | 3             |
| 2019-03-23  | 2             |
| 2019-03-30  | 2             |

# Aggregation - COUNT and COUNT DISTINCT

If we wanted to know how many different products—with unique product IDs—each vendor brought to market during a date range, we could use `COUNT DISTINCT` on the `product_id` field, like so:

```
SELECT
    vendor_id,
    COUNT(DISTINCT product_id) AS different_products_offered
FROM farmers_market.vendor_inventory
WHERE market_date BETWEEN '2019-03-02' AND '2019-03-16'
GROUP BY vendor_id
ORDER BY vendor_id
```

| vendor_id | different_products_offered |
|-----------|----------------------------|
| 1         | 2                          |
| 4         | 1                          |
| 7         | 2                          |
| 8         | 1                          |
| 9         | 3                          |



# Aggregation - Average

What if we also want the **average** original price of a product per vendor?

We can add a line to the preceding query, and use the **AVG()** function:

```
SELECT
    vendor_id,
    COUNT(DISTINCT product_id) AS different_products_offered,
    AVG(original_price) AS average_product_price
FROM farmers_market.vendor_inventory
WHERE market_date BETWEEN '2019-03-02' AND '2019-03-16'
GROUP BY vendor_id
ORDER BY vendor_id
```

| vendor_id | different_products_offered | average_product_price |
|-----------|----------------------------|-----------------------|
| 1         | 2                          | 9.000000              |
| 4         | 1                          | 2.000000              |
| 7         | 2                          | 4.500000              |
| 8         | 1                          | 4.000000              |
| 9         | 3                          | 14.750000             |



# Aggregation - Average

However, we have to think about what we're actually averaging here. Is it fair to call it "average product price" when the underlying table has one row per type of product?

To get an average item price for each vendor's inventory within specified dates, multiply each item's quantity by its price per row, sum these products, and then divide by the total quantity of items per vendor.

```
SELECT
    vendor_id,
    COUNT(DISTINCT product_id) AS different_products_offered,
    SUM(quantity * original_price) AS value_of_inventory,
    SUM(quantity) AS inventory_item_count,
    ROUND(SUM(quantity * original_price) / SUM(quantity), 2) AS
average_item_price
FROM farmers_market.vendor_inventory
WHERE market_date BETWEEN '2019-03-02' AND '2019-03-16'
GROUP BY vendor_id
ORDER BY vendor_id
```

# Aggregation - Average

We also surrounded the calculation with a `ROUND()` function to format the output in dollars:

| vendor_id | different_products_offered | value_of_inventory | inventory_item_count | average_item_price |
|-----------|----------------------------|--------------------|----------------------|--------------------|
| 1         | 2                          | 636.0000           | 68.00                | 9.35               |
| 4         | 1                          | 258.0000           | 129.00               | 2.00               |
| 7         | 2                          | 105.0000           | 27.00                | 3.89               |
| 8         | 1                          | 400.0000           | 100.00               | 4.00               |
| 9         | 3                          | 410.0000           | 30.00                | 13.67              |

# Aggregation - Filtering with HAVING

If you want to filter values after the aggregate functions are applied, you can add a **HAVING** clause to the query. This filters the groups based on the summary values.

Let's filter to vendors who brought at least 100 items to the farmer's market over the specified time period.

```
SELECT
    vendor_id,
    COUNT(DISTINCT product_id) AS different_products_offered,
    SUM(quantity * original_price) AS value_of_inventory,
    SUM(quantity) AS inventory_item_count,
    SUM(quantity * original_price) / SUM(quantity) AS average_item_price
FROM farmers_market.vendor_inventory
WHERE market_date BETWEEN '2019-03-02' AND '2019-03-16'
GROUP BY vendor_id
HAVING inventory_item_count >= 100
ORDER BY vendor_id
```

| vendor_id | different_products_offered | value_of_inventory | inventory_item_count | average_item_price |
|-----------|----------------------------|--------------------|----------------------|--------------------|
| 4         | 1                          | 258.0000           | 129.00               | 2.00000000         |
| 8         | 1                          | 400.0000           | 100.00               | 4.00000000         |

# Aggregation - CASE Statements Inside

We'll use a **CASE** statement to specify which type of item quantities to add together using each **SUM** aggregate function.

First, we'll need to **JOIN** the **customer\_purchases** table to the **product** table to pull in the **product\_qty\_type** column, which only contains the values "unit" and "lbs":

**SELECT**

```
cp.market_date,
cp.vendor_id,
cp.customer_id,
cp.product_id,
cp.quantity,
p.product_name,
p.product_size,
p.product_qty_type
```

```
FROM farmers_market.customer_purchases AS cp
INNER JOIN farmers_market.product AS p
ON cp.product_id = p.product_id
```

| market_date | vendor_id | customer_id | product_id | quantity | product_name         | product_size | product_qty_type |
|-------------|-----------|-------------|------------|----------|----------------------|--------------|------------------|
| 2019-03-02  | 8         | 10          | 4          | 1.00     | Banana Peppers - Jar | 8 oz         | unit             |
| 2019-03-09  | 8         | 12          | 4          | 1.00     | Banana Peppers - Jar | 8 oz         | unit             |
| 2019-03-13  | 8         | 2           | 4          | 2.00     | Banana Peppers - Jar | 8 oz         | unit             |
| 2019-03-16  | 8         | 10          | 4          | 1.00     | Banana Peppers - Jar | 8 oz         | unit             |
| 2019-03-02  | 4         | 2           | 9          | 4.60     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-02  | 4         | 3           | 9          | 8.40     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-02  | 4         | 4           | 9          | 1.40     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-09  | 4         | 4           | 9          | 9.90     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-13  | 4         | 2           | 9          | 4.10     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-16  | 4         | 3           | 9          | 5.50     | Sweet Potatoes       | medium       | lbs              |
| 2019-03-02  | 1         | 1           | 10         | 1.00     | Eggs                 | 1 dozen      | unit             |
| 2019-03-02  | 1         | 1           | 10         | 3.00     | Eggs                 | 1 dozen      | unit             |
| 2019-03-09  | 1         | 1           | 10         | 2.00     | Eggs                 | 1 dozen      | unit             |
| 2019-03-09  | 1         | 4           | 10         | 1.00     | Eggs                 | 1 dozen      | unit             |
| 2019-03-02  | 1         | 1           | 11         | 1.70     | Pork Chops           | 1 lb         | lbs              |
| 2019-03-09  | 1         | 12          | 11         | 0.90     | Pork Chops           | 1 lb         | lbs              |

# Aggregation - CASE Statements Inside

To create columns for quantities sold by **unit**, **pound**, and **other** future units, use **CASE** statements within **SUM** functions to specify which values to add up in each column.

First, we'll review the results with the **CASE** statements included before grouping or using aggregate functions.

```
SELECT
    cp.market_date,
    cp.vendor_id,
    cp.customer_id,
    cp.product_id,
    CASE WHEN product_qty_type = "unit" THEN quantity ELSE 0 END AS
quantity_units,
    CASE WHEN product_qty_type = "lbs" THEN quantity ELSE 0 END AS
quantity_lbs,
    CASE WHEN product_qty_type NOT IN ("unit","lbs") THEN quantity ELSE
0 END AS quantity_other,
    p.product_qty_type
FROM farmers_market.customer_purchases cp
    INNER JOIN farmers_market.product p
        ON cp.product_id = p.product_id
```



# Aggregation - CASE Statements Inside

Notice that the **CASE** statements have separated the quantity values into three different columns, by **product\_qty\_type**.

| market_date | vendor_id | customer_id | product_id | quantity_units | quantity_lbs | quantity_other | product_qty_type |
|-------------|-----------|-------------|------------|----------------|--------------|----------------|------------------|
| 2019-03-02  | 8         | 4           | 4          | 2.00           | 0            | 0              | unit             |
| 2019-03-02  | 8         | 10          | 4          | 1.00           | 0            | 0              | unit             |
| 2019-03-09  | 8         | 12          | 4          | 1.00           | 0            | 0              | unit             |
| 2019-03-13  | 8         | 2           | 4          | 2.00           | 0            | 0              | unit             |
| 2019-03-16  | 8         | 10          | 4          | 1.00           | 0            | 0              | unit             |
| 2019-03-02  | 4         | 2           | 9          | 0              | 4.60         | 0              | lbs              |
| 2019-03-02  | 4         | 3           | 9          | 0              | 8.40         | 0              | lbs              |
| 2019-03-02  | 4         | 4           | 9          | 0              | 1.40         | 0              | lbs              |
| 2019-03-09  | 4         | 4           | 9          | 0              | 9.90         | 0              | lbs              |
| 2019-03-13  | 4         | 2           | 9          | 0              | 4.10         | 0              | lbs              |
| 2019-03-16  | 4         | 3           | 9          | 0              | 5.50         | 0              | lbs              |
| 2019-03-02  | 1         | 1           | 10         | 1.00           | 0            | 0              | unit             |
| 2019-03-02  | 1         | 1           | 10         | 3.00           | 0            | 0              | unit             |
| 2019-03-09  | 1         | 1           | 10         | 2.00           | 0            | 0              | unit             |
| 2019-03-09  | 1         | 4           | 10         | 1.00           | 0            | 0              | unit             |
| 2019-03-02  | 1         | 1           | 11         | 0              | 1.70         | 0              | lbs              |
| 2019-03-09  | 1         | 12          | 11         | 0              | 0.90         | 0              | lbs              |



# Aggregation - CASE Statements Inside

Now we can add the **SUM** functions around each **CASE** statement to add up these values per market date per customer, as defined in the **GROUP BY** clause.

```
SELECT
    cp.market_date,
    cp.customer_id,
    SUM(CASE WHEN product_qty_type = "unit" THEN quantity ELSE 0 END) AS
    qty_units_purchased,
    SUM(CASE WHEN product_qty_type = "lbs" THEN quantity ELSE 0 END) AS
    qty_lbs_purchased,
    SUM(CASE WHEN product_qty_type NOT IN ("unit","lbs") THEN quantity
    ELSE 0 END) AS qty_other_purchased
FROM farmers_market.customer_purchases cp
    INNER JOIN farmers_market.product p
        ON cp.product_id = p.product_id
GROUP BY market_date, customer_id
ORDER BY market_date, customer_id
```

| market_date | customer_id | qty_units_purchased | qty_lbs_purchased | qty_other_purchased |
|-------------|-------------|---------------------|-------------------|---------------------|
| 2019-03-02  | 1           | 4.00                | 1.70              | 0.00                |
| 2019-03-02  | 2           | 0.00                | 4.60              | 0.00                |
| 2019-03-02  | 3           | 0.00                | 8.40              | 0.00                |
| 2019-03-02  | 4           | 2.00                | 1.40              | 0.00                |
| 2019-03-02  | 10          | 1.00                | 0.00              | 0.00                |
| 2019-03-09  | 1           | 2.00                | 2.20              | 0.00                |
| 2019-03-09  | 4           | 3.00                | 10.20             | 0.00                |
| 2019-03-09  | 7           | 2.00                | 0.00              | 0.00                |
| 2019-03-09  | 12          | 3.00                | 0.90              | 0.00                |
| 2019-03-13  | 2           | 2.00                | 4.10              | 0.00                |
| 2019-03-16  | 3           | 0.00                | 5.50              | 0.00                |
| 2019-03-16  | 10          | 1.00                | 0.00              | 0.00                |
| 2019-03-20  | 1           | 0.00                | 3.10              | 0.00                |
| 2019-03-20  | 7           | 3.00                | 0.00              | 0.00                |
| 2019-03-23  | 4           | 3.00                | 2.40              | 0.00                |

# Exercise

1. Write a query that determines how many times each vendor has rented a booth at the farmer's market. In other words, count the vendor booth assignments per `vendor_id`.
2. We asked earlier "When is each type of fresh fruit or vegetable in season, locally?" Write a query that displays the product category name, product name, earliest date available, and latest date available for every product in the "Fresh Fruits & Vegetables" product category.
3. The Farmer's Market Customer Appreciation Committee wants to give a bumper sticker to everyone who has ever spent more than \$50 at the market. Write a query that generates a list of customers for them to give stickers to, sorted by last name, then first name. (**HINT**: This query requires you to join two tables, use an aggregate function, and use the HAVING keyword.)

# Q&A

## Questions and answers

# Thanks!